

# Programming Project #5: Video Stitching and Processing

## CS445: Computational Photography - Fall 2019

### Part I: Stitch two key frames

This involves:

1. compute homography  $H$  between two frames;
2. project each frame onto the same surface;
3. blend the surfaces.



Check that your homography is correct by plotting four points that form a square in frame 270 and their projections in each image, like this:



```
In [1]: 1 import cv2
        2 import numpy as np
        3 from numpy.linalg import svd, inv, solve, lstsq
        4 import scipy
        5 from scipy.sparse.linalg import lsqr
        6 from scipy.sparse import csr_matrix
        7 %matplotlib inline
        8 from matplotlib import pyplot as plt
```

```
In [2]: 1 # images location
        2 im1 = './images/input/frames/f0001.jpg'
        3 im2 = './images/input/frames/f0270.jpg'
        4
        5 # Load an color image in grayscale
        6 im1 = cv2.imread(im1)
        7 im2 = cv2.imread(im2)
```

```

In [3]: 1 def auto_homography(Ia,Ib, homography_func=None, normalization_func=None):
2         '''
3         Computes a homography that maps points from Ia to Ib
4
5         Input: Ia and Ib are images
6         Output: H is the homography
7
8         '''
9         if Ia.dtype == 'float32' and Ib.dtype == 'float32':
10            Ia = (Ia*255).astype(np.uint8)
11            Ib = (Ib*255).astype(np.uint8)
12
13            Ia_gray = cv2.cvtColor(Ia,cv2.COLOR_BGR2GRAY)
14            Ib_gray = cv2.cvtColor(Ib,cv2.COLOR_BGR2GRAY)
15
16            # Initiate SIFT detector
17            sift = cv2.xfeatures2d.SIFT_create()
18
19            # find the keypoints and descriptors with SIFT
20            kp_a, des_a = sift.detectAndCompute(Ia_gray,None)
21            kp_b, des_b = sift.detectAndCompute(Ib_gray,None)
22
23            # BFMatcher with default params
24            bf = cv2.BFMatcher()
25            matches = bf.knnMatch(des_a,des_b, k=2)
26
27            # Apply ratio test
28            good = []
29            for m,n in matches:
30                if m.distance < 0.75*n.distance:
31                    good.append(m)
32
33            numMatches = int(len(good))
34
35            matches = good
36
37            # Xa and Xb are 3xN matrices that contain homogeneous coordinates for the
38            # matching points for each image
39            Xa = np.ones((3,numMatches))
40            Xb = np.ones((3,numMatches))
41
42            for idx, match_i in enumerate(matches):
43                Xa[:,idx][0:2] = kp_a[match_i.queryIdx].pt
44                Xb[:,idx][0:2] = kp_b[match_i.trainIdx].pt
45
46            ## RANSAC
47            niter = 1000
48            best_score = 0
49
50            for t in range(niter):
51                # estimate homography
52                subset = np.random.choice(numMatches, 4, replace=False)
53                pts1 = Xa[:,subset]
54                pts2 = Xb[:,subset]
55
56                H_t = homography_func(pts1, pts2, normalization_func) # edit helper c

```

```

57     #H_t = computeHomography(pts1, pts2)
58
59     # score homography
60     Xb_ = np.dot(H_t, Xa) # project points from first image to second using H_t
61     du = Xb_[0,:]/Xb_[2,:] - Xb[0,:]/Xb[2,:]
62     dv = Xb_[1,:]/Xb_[2,:] - Xb[1,:]/Xb[2,:]
63
64     ok_t = np.sqrt(du**2 + dv**2) < 1 # you may need to play with this threshold
65     score_t = sum(ok_t)
66
67     if score_t > best_score:
68         print("get a better score")
69         best_score = score_t
70         H = H_t
71         in_idx = ok_t
72
73     print('best score: {:.02f}'.format(best_score))
74
75     # Optionally, you may want to re-estimate H based on inliers
76
77     return H

```

In [4]:

```

1  from scipy import linalg
2
3  def get_A(src,dst):
4      eye = np.eye(3)
5      zeros = np.zeros([3,3])
6      first_half = np.concatenate((eye*-1,zeros,dst[0]*eye), axis=-1)
7      second_half = np.concatenate((zeros,eye*-1,dst[1]*eye), axis=-1)
8      #print(src.T.dot(first_half).shape, src.T.dot(second_half).shape)
9      return src.T.dot(first_half),src.T.dot(second_half)
10
11 def computeHomography(pts1, pts2, normalization_func=None):
12     '''
13     Compute homography that maps from pts1 to pts2 using least squares solver
14
15     Input: pts1 and pts2 are 3xN matrices for N points in homogeneous
16     coordinates.
17
18     Output: H is a 3x3 matrix, such that pts2~H*pts1 pts1.T * H^T = pts2.T
19     '''
20     N = pts1.shape[1]
21     l = []
22     for i in range(N):
23         a,b = get_A(pts1[:,i],pts2[:,i])
24         l.append(a)
25         l.append(b)
26     #print(l)
27     A = np.asarray(l)
28     #print(A.shape)
29     #print(A)
30     U,s,Vt = np.linalg.svd(A)
31     #print(Vt.shape)
32     H = Vt[-1,:].reshape([3,3])
33     return H
34

```

```
In [5]: 1 H = auto_homography(im1,im2, computeHomography)
```

```
get a better score  
get a better score  
get a better score  
get a better score  
get a better score  
get a better score  
get a better score  
best score: 159.000000
```

In [6]:

```

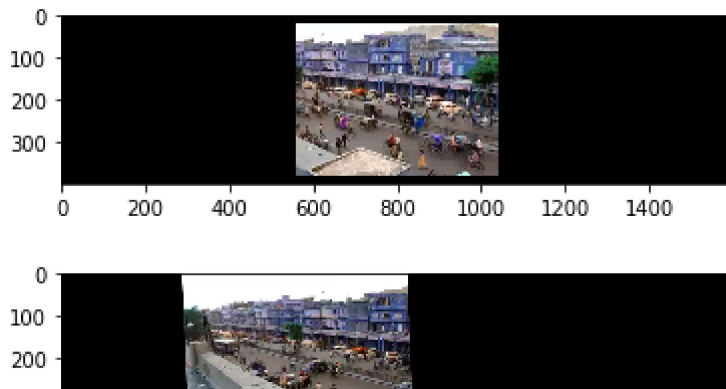
1
2 output_width, output_height = (1600, 400)
3
4 corners = np.array([np.array([[0,0],[im1.shape[1]-1,0],[0,im1.shape[0]-1],[im1.shape[1]-1,im1.shape[0]-1]]),
5 center = np.array([np.array([output_width//2-im1.shape[1]//2, output_height//2,
6 [output_width//2+im1.shape[1]//2, output_height//2,
7 [output_width//2-im1.shape[1]//2, output_height//2,
8 [output_width//2+im1.shape[1]//2, output_height//2,
9 #print(corners)
10 H_t = cv2.getPerspectiveTransform(corners,center)
11 print(H_t)
12 shift = np.array([[1,0,800],[0,1,0],[0,0,1]]).astype(np.float32)
13 shift2 = np.array([[0,0,800],[0,0,0],[0,0,0]]).astype(np.float32)
14 img_warped = cv2.warpPerspective(im2, H_t, (output_width, output_height)).ast
15 plt.imshow(img_warped/255.)
16 plt.show()
17
18 img_warped2 = cv2.warpPerspective(im1, H_t.dot(H), (output_width, output_heig
19 mask = np.where(img_warped>0, 1, 0).astype(np.float32)
20 plt.imshow(img_warped2/255.)
21 plt.show()
22
23 print(np.max(img_warped))
24 print("mask")
25 plt.imshow(mask.astype(np.float32))
26 plt.show()
27 print("background")
28 plt.imshow(img_warped2.astype(np.float32)/255.)
29 plt.show()
30 print("foreground")
31 plt.imshow(img_warped.astype(np.float32)/255.)
32 plt.show()
33 print(np.sum(mask,axis=-1).shape)
34 #img_warped3 = laplacian_blend(img_warped, img_warped2, np.sum(mask,axis=-1))
35 #img_warped3 = mask*img_warped+(1.-mask)*img_warped2
36 #alpha blend
37 img_warped3 = mask*img_warped+(1.-mask)*img_warped2
38
39 result = cv2.cvtColor(img_warped3,cv2.COLOR_BGR2RGB)
40
41 print(H)
42 plt.imshow(result/255.)
43 plt.show()
44 print(np.max(result))
45 plt.imsave("trial1.jpg",np.clip(result/255.,0.,1.))
46 del result
47 del img_warped, img_warped2, img_warped3

```

```

[[ 1.00208768e+00 -6.66133815e-16  5.60000000e+02]
 [-1.33226763e-15  1.00278552e+00  2.00000000e+01]
 [-1.73472348e-18  4.33680869e-19  1.00000000e+00]]

```



## Part II: Panorama using five key frames

In this part you will produce a panorama using five key frames. Let's determine frames [90, 270, 450, 630, 810] as key frames. The goal is to map all the five frames onto the plane corresponding to frame 450 (that we also call the *reference frame*). For the frames 270 and 630 you can follow the instructions in part 1.



Mapping frame 90 to frame 450 is difficult because they share very little area. Therefore you need to perform a two stage mapping by using frame 270 as a guide. Compute one projection from 90 to 270 and one from 270 to 450 and multiply the two matrices. This produces a projection from 90 to 450 even though these frames have very little area in common

```
In [7]: 1 import cv2
        2 import numpy as np
```

```

In [8]: 1 master_frames =[90, 270, 450, 630, 810]
2 output_width, output_height = 1600,500
3 img_list = []
4 img_path_list=['./images/input/frames/f%04d.jpg'%k for k in master_frames]
5 print(img_path_list)
6 img_list = [cv2.imread(t) for t in img_path_list]
7
8 reference_frame = 450
9 reference_idx = master_frames.index(reference_frame)
10 H_list = []
11 corners = np.array([np.array([0,0],[im1.shape[1]-1,0],[0,im1.shape[0]-1],[im
12 center = np.array([np.array([output_width//2-im1.shape[1]//2, output_height/
13 [output_width//2+im1.shape[1]//2, output_height/
14 [output_width//2-im1.shape[1]//2, output_height/
15 [output_width//2+im1.shape[1]//2, output_height/
16 H_t = cv2.getPerspectiveTransform(corners,center)
17 shift = np.array([1,0,output_width//2-img_list[reference_idx].shape[1]//2],[
18 shift2 = np.array([0,0,output_width//2-img_list[reference_idx].shape[1]//2],
19 for i in range(reference_idx):
20     H_list.append(auto_homography(img_list[i],img_list[i+1], computeHomograph
21 for i in range(reference_idx+1, len(master_frames)):
22     H_list.append(auto_homography(img_list[i],img_list[i-1], computeHomograph
23
24 img_warped = []
25 mask = []
26 for i in range(reference_idx):
27     #print("image: ",i)
28     H = H_list[i]
29     #print(H)
30     for j in range(reference_idx):
31         if j>i:
32             print(H_list[j])
33             H = np.dot(H_list[j],H)
34     #print("H: ")
35     #print(H)
36     img_warped.append(cv2.warpPerspective(img_list[i], H_t.dot(H), (output_wi
37     mask.append(np.where(img_warped[-1]>0,1.,0.))
38
39 img_warped.append(cv2.warpPerspective(img_list[reference_idx], H_t, (output_w
40 mask.append(np.where(img_warped[-1]>0,1.,0.))
41
42 for i in range(reference_idx+1, len(master_frames)):
43     H = H_list[i-1]
44     for j in range(i-2,0,-1):
45         if j>=reference_idx:
46             H = np.dot(H_list[j],H)
47     img_warped.append(cv2.warpPerspective(img_list[i], H_t.dot(H), (output_wi
48     mask.append(np.where(img_warped[-1]>0,1.,0.))
49
50
51 result = np.zeros((output_height,output_width,3))
52 for i in range(len(master_frames)-1,-1,-1):
53     #print("result: ",result.shape)
54     #print("mask: ",mask[i].shape)
55     #print("img: ",img_list[i].shape)
56     result = result*(1-mask[i])+img_warped[i]

```



```

57
58 result = cv2.cvtColor(np.clip(result.astype(np.uint8),0,255),cv2.COLOR_BGR2RG
59 plt.imshow(np.clip(result/255.,0.,1.))
60 plt.show()
61 plt.imsave("panoramic.jpg",result)
62 #img_warped3 = laplacian_blend(img_warped, img_warped2, np.sum(mask,axis=-1))
63 #img_warped3 = mask*img_warped+(1.-mask)*img_warped2
64 #alpha blend
65 #img_warped3 = mask*img_warped+(1.-mask)*img_warped2
66
67 #result = cv2.cvtColor(img_warped3,cv2.COLOR_BGR2RGB)
68 del img_list, result, img_warped

```

```

['./images/input/frames/f0090.jpg', './images/input/frames/f0270.jpg', './ima
ges/input/frames/f0450.jpg', './images/input/frames/f0630.jpg', './images/inp
ut/frames/f0810.jpg']

```

```

get a better score
get a better score
get a better score
get a better score
best score: 210.000000
get a better score
get a better score
get a better score
get a better score
get a better score
get a better score
best score: 158.000000
get a better score
get a better score
get a better score
get a better score
get a better score

```

### Part 3: Map the video to the reference plane

```

In [9]: 1 import os
        2 import cv2
        3 import numpy as np
        4 import matplotlib.pyplot as plt
        5 from math import floor
        6
        7 import utils

```

```

In [10]: 1 dir_frames = 'images/input/frames'
        2 filenames = []
        3 filesinfo = os.scandir(dir_frames)

```

```

In [11]: 1 filenames = [f.path for f in filesinfo if f.name.endswith(".jpg")]
        2 filenames.sort(key=lambda f: int(''.join(filter(str.isdigit, f))))

```

```
In [12]: 1 frameCount = len(filenamees)
          2 frameHeight, frameWidth, frameChannels = cv2.imread(filenamees[0]).shape
          3 frames = np.zeros((frameCount, frameHeight, frameWidth, frameChannels),dtype=

In [13]: 1 for idx, file_i in enumerate(filenamees):
          2     frames[idx] = cv2.cvtColor(cv2.imread(file_i), cv2.COLOR_BGR2RGB) / 255.0
```

In [28]:

```

1  ## Example usage of utils.projectImage
2  H = 700
3  W = 2000
4  #pastHomographies = np.zeros((len(filenames),len(filenames), 3, 3),dtype=np.f
5  originTranslations = np.zeros((len(filenames), 2), dtype=np.float32)
6  originTranslations[:,0]-=1000-int(frames[0].shape[1]/2)
7  originTranslations[:,1]-=350-int(frames[0].shape[0]/2)
8  #print(originTranslations[451:])
9  #originTranslations[451:,0]+=1000-int(frames[0].shape[1]/2)#1000-int(frames[0
10 #originTranslations[451:,1]+=350-int(frames[0].shape[0]/2)#350-int(frames[0].
11 referenceFrameIndex = 450
12
13 for i in range(450):#len(frames)):
14     sourceFrameIndex = int(i)
15     print(i)
16     #print(len(frames))
17     #print(frames[0].shape)
18     projectedSource, pastHomographies, originTranslations = utils.projectImag
19                                                         pastHomograph
20                                                         yrange=H, ove
21                                                         numKeyframes=
22                                                         auto_H_func=a
23     #plt.imshow(projectedSource)
24     #plt.show()
25     #print(past_trans1.shape)
26     name = cv2.imwrite('aligned_frames/a{:04d}.jpg'.format(i), cv2.cvtColor(p
27     #trans_map[i,:,:]=past_trans1[:,:]
28     ...
29
30
31 for i in range(frameCount-1,450,-1):#len(frames)):
32     sourceFrameIndex = int(i)
33     print(i)
34     #print(len(frames))
35     #print(frames[0].shape)
36     projectedSource, pastHomographies, originTranslations = utils.projectImag
37                                                         pastHomograph
38                                                         yrange=H, ove
39                                                         numKeyframes=
40                                                         auto_H_func=a
41     #plt.imshow(projectedSource)
42     #plt.show()
43     name = cv2.imwrite('aligned_frames/a{:04d}.jpg'.format(i), cv2.cvtColor(p
44     #trans_map[i,:,:]=past_trans2[:,:]
45     ...
46 np.save("trans",pastHomographies)
47     #utils.imageFolder2mpeg('aligned_frames', fps=30)

```

```

0
Overlap:169200
Error:0.000493949100987447
Finding better homography...
Overlap:135386
Error:0.0005839388238521869
Overlap:148966

```

```
Error:0.00018655811193993158
Overlap:67335
Error:0.0007513987659576506
get a better score
get a better score
get a better score
get a better score
get a better score
best score: 154.000000
Overlap:106438
Error:0.00027355704899158716
Overlap:5400
```

```
In [29]: 1 utils.imageFolder2mpeg('aligned_frames', fps=30)
```

## Part 4: Create background panorama

In this part you will remove moving objects from the video and create a background panorama that should incorporate pixels from all the frames.

In the video you produced in **part 3** each pixel appears in several frames. You need to estimate which of the many colors correspond to the background. We take advantage of the fact that the background color is fixed while the foreground color changes frequently (because foreground moves).



For each pixel in the sequence of **part 3**, determine all valid colors (colors that come from all frames that overlap that pixel). You can experiment with different methods for determining the background color of each pixel, as discussed in class. Perform the same procedure for all pixels and generate output. The output should be a completed panorama showing only pixels of background or non-moving objects.

```
In [30]: 1 import os
          2 import cv2
          3 import numpy as np
          4 import matplotlib.pyplot as plt
```

In [34]:

```

1  ## Example usage of utils.projectImage
2  H = 512
3  W = 1632
4  mean = np.zeros((H,W,3))
5  count = np.zeros((H,W,3))
6  print(mean.shape)
7
8  filenames = [f.path for f in os.scandir('./aligned_frames/') if f.name.endswi
9  #print(filenames)
10 filenames.sort(key=lambda f: int(f[len('./aligned_frames/')+1:len('./aligned_
11 print(filenames)
12 frame_list = [plt.imread(f) for f in filenames]
13 num_frames = len(frame_list)

```

(512, 1632, 3)

['./aligned\_frames/a0000.jpg', './aligned\_frames/a0001.jpg', './aligned\_frame  
s/a0002.jpg', './aligned\_frames/a0003.jpg', './aligned\_frames/a0004.jpg', './  
aligned\_frames/a0005.jpg', './aligned\_frames/a0006.jpg', './aligned\_frames/a0  
007.jpg', './aligned\_frames/a0008.jpg', './aligned\_frames/a0009.jpg', './alig  
ned\_frames/a0010.jpg', './aligned\_frames/a0011.jpg', './aligned\_frames/a0012.  
jpg', './aligned\_frames/a0013.jpg', './aligned\_frames/a0014.jpg', './aligned\_  
frames/a0015.jpg', './aligned\_frames/a0016.jpg', './aligned\_frames/a0017.jp  
g', './aligned\_frames/a0018.jpg', './aligned\_frames/a0019.jpg', './aligned\_fr  
ames/a0020.jpg', './aligned\_frames/a0021.jpg', './aligned\_frames/a0022.jpg',  
'./aligned\_frames/a0023.jpg', './aligned\_frames/a0024.jpg', './aligned\_frame  
s/a0025.jpg', './aligned\_frames/a0026.jpg', './aligned\_frames/a0027.jpg', './  
aligned\_frames/a0028.jpg', './aligned\_frames/a0029.jpg', './aligned\_frames/a0  
030.jpg', './aligned\_frames/a0031.jpg', './aligned\_frames/a0032.jpg', './alig  
ned\_frames/a0033.jpg', './aligned\_frames/a0034.jpg', './aligned\_frames/a0035.  
jpg', './aligned\_frames/a0036.jpg', './aligned\_frames/a0037.jpg', './aligned\_  
frames/a0038.jpg', './aligned\_frames/a0039.jpg', './aligned\_frames/a0040.jp  
g', './aligned\_frames/a0041.jpg', './aligned\_frames/a0042.jpg', './aligned\_fr  
ames/a0043.jpg', './aligned\_frames/a0044.jpg', './aligned\_frames/a0045.jpg',  
'./aligned\_frames/a0046.jpg', './aligned\_frames/a0047.jpg', './aligned\_frame

In [35]:

```

1  #del frames
2  for t in range(len(frame_list)):
3      if frame_list[t].shape!=(H,W,3):
4          print(t,frame_list[t].shape)

```

```

In [36]: 1 frame_list = np.concatenate(frame_list,axis=-1).reshape((H,W,frameCount,3))
2
3 frame_color_R = np.zeros((H,W,frameCount))
4 frame_color_R += frame_list[:, :, :, 0]
5 frame_color_G = np.zeros((H,W,frameCount))
6 frame_color_G += frame_list[:, :, :, 1]
7 frame_color_B = np.zeros((H,W,frameCount))
8 frame_color_B += frame_list[:, :, :, 2]
9 del frame_list
10 R = np.zeros((H,W))
11 frame_color_R = np.sort(frame_color_R,axis=-1)
12 for y in range(H):
13     print(y)
14     for x in range(W):
15         count = -1
16         while frame_color_R[y,x,count]:
17             count-=1;
18         R[y,x] = np.median(frame_color_R[y,x,count:])
19 del frame_color_R
20
21 G = np.zeros((H,W))
22 frame_color_G = np.sort(frame_color_G,axis=-1)
23 for y in range(H):
24     print(y)
25     for x in range(W):
26         count = -1
27         while frame_color_G[y,x,count]:
28             count-=1;
29         G[y,x] = np.median(frame_color_G[y,x,count:])
30 del frame_color_G
31
32 B = np.zeros((H,W))
33 frame_color_B = np.sort(frame_color_B,axis=-1)
34 for y in range(H):
35     print(y)
36     for x in range(W):
37         count = -1
38         while frame_color_B[y,x,count]:
39             count-=1;
40         B[y,x] = np.median(frame_color_B[y,x,count:])
41 del frame_color_B
42
43 result = cv2.merge([R,G,B])
44
45
46 #plt.imshow(frame_list[:, :, 0, :].reshape((H,W,3)))
47 # =
48 #color_map =

```

0  
1  
2  
3  
4  
5  
6

7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18

In [37]:

```
1  
2  
3 print(result.shape)  
4 print(np.max(result))  
5 plt.imshow(result.astype(np.uint8))  
6 plt.show()  
7  
8 plt.imsave("median_background.jpg", (result).astype(np.uint8))
```

(512, 1632, 3)  
255.0



## Part 5: Create background movie

Map the background panorama to the movie coordinates. For each frame of the movie, say frame 1, you need to estimate a projection from the panorama to frame 1. Note, you should be able to re-use the homographies that you estimated in **Part 3**. Perform this for all frames and generate a movie that looks like the input movie but shows only background pixels. All moving objects that belong to the foreground must be removed.

In [38]:

```
1 import os  
2 import cv2  
3 import numpy as np  
4 pastHomographies = np.load("trans.npy")  
5 result = plt.imread("median_background.jpg")
```

In [39]:

```

1  ## Example usage of utils.projectImage
2
3  mean = np.zeros((H,W,3))
4  count = np.zeros((H,W,3))
5  print(mean.shape)
6
7  filenames = [f.path for f in os.scandir('./aligned_frames/') if f.name.endswi
8  #print(filenames)
9  filenames.sort(key=lambda f: int(f[len('./aligned_frames/')+1:len('./aligned_
10 #print(filenames)
11 frame_list = [plt.imread(f) for f in filenames]
12 for i in range(len(frame_list)):
13     frame = frame_list[i]
14     #plt.imshow(frame)
15     #plt.show()
16     #foreground_mask = np.where(np.abs(frame-median)>thredhold, 1, 0)
17     background_mask = np.where(frame,1.,0.)
18     #plt.imshow(background_mask)
19     #plt.show()
20     #print(background_mask.shape)
21     print(i)
22     #print(pastHomographies[i,450])
23     H = np.linalg.inv(pastHomographies[i,450])
24     background_frame = result*background_mask
25     warped = cv2.warpPerspective(background_frame, H, (480,360)).astype(np.ui
26     #plt.imshow(background_frame/255.)
27     #plt.show()
28     name = plt.imsave('aligned_background_frames/a{:04d}.jpg'.format(i), warp
29

```

(512, 1632, 3)

```

0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18

```

In [41]:

```

1  utils.imageFolder2mpeg('aligned_background_frames', './output_background_video

```

## Part 6: Create foreground movie



In the background video, moving objects are removed. In each frame, those pixels that are different enough than the background color are considered foreground. For each frame determine foreground pixels and generate a movie that only includes foreground pixels.

```
In [44]: 1 import os
2 import cv2
3 import numpy as np
4
5 ## Example usage of utils.projectImage
6 print(mean.shape)
7
8 filenames = [f.path for f in os.scandir('./aligned_frames/') if f.name.endswith
9 #print(filenames)
10 filenames.sort(key=lambda f: int(f[len('./aligned_frames/')+1:len('./aligned_
11 #print(filenames)
12 frame_list = [plt.imread(f) for f in filenames]
13 for i in range(len(frame_list)):
14     frame = frame_list[i]
15     print(i)
16     background_mask = np.where(frame,1.,0.)
17     background_frame = result*background_mask
18
19     dist = np.sum((frame-background_frame)**2,axis=-1)
20     #print(dist.shape)
21     dist = cv2.merge([dist,dist,dist])
22     foreground_frame_mask = np.where(dist>2000, 1, 0)
23     #print(foreground_frame_mask.shape)
24     foreground_frame = frame*foreground_frame_mask
25     H = np.linalg.inv(pastHomographies[i,450])
26     warped = cv2.warpPerspective(foreground_frame.astype(np.uint8), H, (480,3
27     name = plt.imsave('aligned_foreground_frames/a{:04d}.jpg'.format(i), warp
28
```

```
(512, 1632, 3)
```

```
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
```

```
In [45]: 1 utils.imageFolder2mpeg('aligned_foreground_frames','./output_foreground_video
```

# Bells and whistles

In [ ]:

1